

# How to Make a Proceedings Paper Submission

## Anonymous CogSci submission

### Abstract

Include no author information in the initial submission, to facilitate blind review. The abstract should be one paragraph, indented 1/8 inch on both sides, in 9-point font with single spacing. The heading 'Abstract' should be 10-point, bold, centered, with one line of space below it. This one-paragraph abstract section is required only for standard six page proceedings papers. Following the abstract should be a blank line, followed by the header 'Keywords' and a list of descriptive keywords separated by semicolons, all in 9-point font, as shown below.

**Keywords:** Add your choice of indexing terms or keywords; kindly use a semi-colon; between each term.

### Introduction

Whether to keep looking at a current target of attention is one of the most fundamental decisions we make, whether we are trying to find our way in a busy street or swiping through TikTok. Even young infants are tasked with making the decision on selecting what to look at and for how long. To look or not to look, this decision that infants make constantly has provided developmental researchers an opportunity to investigate infants' mental world. Through the use of looking time paradigms, researchers are able to make inferences about infants' learning and mental representations based on changes in looking time (CITE, CITE, CITE). In a typical experiment, infants increasingly decrease their looking duration upon seeing repeated stimulus (i.e. habituation). When habituated, infants regain their interests when seeing a novel stimulus (i.e. dishabituation). While both phenomena are well-documented, the factors that influence these looking time trajectories remain relatively underexplored. A better understanding of what shapes habituation and dishabituation is critical given their methodological and theoretical significance. The rise and fall in looking time is not only central to understanding infants' mental representation, but also shed light on principles that guide information-seeking behavior in general.

Classical theory of infant looking behavior suggests three factors are crucial to habituation and dishabituation: complexity, familiarization time, and infants' age (Hunter & Ames, 1988). More perceptually complex stimuli take longer time for infants to habituate. Longer familiarization time to one stimulus would make infants more likely to dishabituate to another stimulus. The older infants are, the more efficient they are at information processing, and the faster

they are to habituate when other factors are controlled for. Together, these three factors determine how infants' looking time changes during an experiment. Although Hunter & Ames (1988) is influential, the evidence for the theory is weak, with some studies showing mixed results (CITE meta analysis). Furthermore, this verbal theory lacks quantitative details, and therefore unlikely to offer precise predictions on looking time changes based on the different factors.

In contrast to verbal theory, computational models offer quantitative predictions. More recent work has linked infants' looking behaviors with a range of information theoretic measurements derived from models. In pioneering work, KPA (CITE) developed a paradigm in which infants are shown sequences of events. Infants' look-away probabilities toward the stimuli are compared with surprisal, a measure of information content, derived from a rational learner model that keeps track of the probabilities of each event. The study shows that infants looking behaviors can be predicted by surprisal. In particular, they pay most attention to event sequences that are neither too high nor too low in surprisal. A recent study with a similar paradigm provides an alternative linking hypothesis. In Poli et al (2020), another information theoretic measurement, Kullback-Leibler divergence, is shown to outperform surprisal in predicting infants' looking time. These attempts on connecting information theoretic measurements to infants' looking time resonate with the emerging literature on curiosity in developmental robotics and reinforcement learning (CITE, CITE, CITE). Curiosity-driven artificial agents' exploratory behaviors are guided by optimizing Expected Information Gain (EIG) (CITE, CITE), a measurement that has been shown to be related to information-seeking in human children and adults as well (CITE).

However, there are several limitations to the existing models. First, the current models did not capture the noisy nature of perceptual learning (CITE noisy perception?). The rational learner models were assumed to acquire perfect representation of each event in the sequence (CITE model). This assumption leads to the second limitation: the lack of explanation in why a learner would choose to learn a stimulus in the first place. Both surprisal and KL-divergence have been presented as potential explanations of infants' looking behaviors, yet neither of the measurements is mechanistically linked to the models' behaviors. They are descriptive in nature, derived

from models that track the probabilities of the events. Finally, the behavioral data that the models were evaluated with came from experimental paradigms that were not representatives of infant looking time paradigms. The key phenomena, habituation and dishabituation, were not captured. The extent to which we can extrapolate current behavior-model fits to understand changes in looking time in a typical looking time experiment remains limited.

Here we present a series of models that can explain patterns in looking time. Our Goal is to provide a unifying quantitative account of looking behaviors as arising from optimal decision-making over noisy perceptual representations (CITE C & G; drif diffusion). We begin by instantiating a version of prior learning models in an independent-trial format (where individual stimuli are learned, not sequences of events). We then develop a second model that addresses weakness in previous work by a) assume the model is accumulating noisy samples from the stimulus, and b) assume the model is choosing what to look at depending on the linking hypotheses (surprisal, KL-divergence, and EIG). Finally, we evaluate our model with adult looking time data collected from a paradigm that captures habituation, dishabituation, and complexity effect.

## Models

**Discrete Time Model** We formalized the learning problem that participants face in our experiments as a form of Bayesian concept learning (Goodman, Tenenbaum, Feldman, & Griffiths, 2008; Tenenbaum, 1999), represented graphically in Fig. X. The goal is to learn a concept  $\theta$ , which is a set of probabilities for independent binary features  $\theta_{1,2,\dots,n}$ , where  $n$  is the number of features. Over the course of a block, the learner receives information about  $\theta$  by observing exemplars  $y$ : instantiations of  $\theta$ , where each feature  $y_{1,2,\dots,n}$  is either on or off. Each feature  $\theta_i$  and its corresponding exemplar  $y_i$  form a Beta-Bernoulli process:

$$p(\theta_i) \sim \text{Beta}(\alpha_i, \beta_i) \quad (1)$$

$$p(y_i|\theta_i) \sim \text{Bernoulli}(\theta_i) \quad (2)$$

Since the features are independent, this relationship holds for the entire concept  $\theta$ . In previous work, two information-theoretic quantities, surprisal and Kullback-Leibler (KL) divergence, resulting from the stimulus were shown to be linked to looking behavior (Kidd, Piantadosi, & Aslin, 2012; Poli, Serino, Mars, & Hunnius, 2020). Surprisal, calculated as  $-\log(p(y|\theta))$ , intuitively refers to how surprising a stimulus  $y$  is given the model's beliefs about  $\theta$  - the intuition that surprising events should result in longer looking times has served as a foundational assumption in developmental psychology. KL-divergence measures how much a model needs to change to accommodate a new stimulus  $y$ , and describes a distance between the model before and after an observation, in is defined in our case as  $\sum_{x \in X} p(\theta = x|y) \frac{p(\theta=x|y)}{p(\theta=x)}$ . If an observation causes a large change, we speculated that a proportionally long looking time is necessary to integrate the new

information.

**Continuous Time Model** However, to model the time course of attention, we did not want to assume that stimuli are encoded perfectly and instantaneously. Instead, we suggest that participants gather repeated noisy samples  $\bar{z}$  from the exemplars. For any sample  $z$  from an exemplar  $y$  there is a small probability  $\epsilon$  to misperceive the feature as off when it was actually on, and vice versa. Therefore, by making noisy observations  $\bar{z}$ , the learner obtains information about the true identity of the exemplar  $y$ , and by extension, about the concept  $\theta$ . By Bayes' rule:

$$P(\theta|\bar{z}) = p(\bar{z}|y)p(y|\theta)p(\theta)/p(\bar{z}) \quad (3)$$

where  $p(\bar{z}|y_i)$  is fully described by  $\epsilon$ , and  $p(y|\theta)$  by Bernoulli processes as in Eq. 2.




Figure 1: Graphical representation of our model. Circles indicate random variables. The squares indicate fixed model parameters.

**Sampling** The formulation of the model in continuous time allows us to do two things: First, we can explicitly model the learner's decision on when to stop sampling by asking the model to decide, after every sample  $z$ , whether it wants to continue sampling from the same stimulus or not. This is in contrast to the discrete time models presented here and in previous work (Kidd et al., 2012; Poli et al., 2020), where we can only link information-theoretic measures to looking data, but not provide a mechanism for how these measures could control moment-to-moment sampling decisions. Second, a consequence of making a decision at every time step is that we can study the behavior of another information-theoretic measure: the expected information gain (EIG). EIG is commonly used in rational analyses of information-seeking behavior - that is to assess whether information-seeking is optimal with respect to the learning task (Markant & Gureckis, 2012; Oaksford & Chater, 1994). Importantly, EIG is a forward-looking measure that considers the potential for learning from the next sample. Since discrete time models operate on the level of a whole stimulus, rather than individual samples, EIG would look forward to the next stimulus in these models,

rather than the next sample, and therefore not be able to capture the decision of whether to keep looking. EIG to describe looking time is therefore only possible in the continuous time models.

We compute EIG by weighing the information gain from each possible next observation by the probability of that observation. We defined information gain as the KL-divergence between the hypothetical posterior after observing a sample  $z_{t+1}$  and the current posterior:

$$EIG(z_{t+1}) = \sum_{z_{t+1} \in [0,1]} p(z_{t+1}|\theta_t) * KL(\theta_{t+1}, p(\theta_t)) \quad (4)$$

Finally, to get actual sampling behavior from the model, it has to convert EIG into a binary decision about whether continue looking at the current sample, or to advance to the next trial. The model does so using a luce choice between the EIG from the next sample and a constant EIG from looking away.

$$p(\text{look}) = \frac{EIG(z_{t+1})}{EIG(z_{t+1}) + EIG(\text{world})} \quad (5)$$

We also studied the behavior of the model when replacing EIG with continuous time versions of the other linking hypotheses, surprisal and KL-divergence between the posterior  $p(\theta_t)$  and the prior  $p(\theta_{t-1})$ .

## Experiment

### Methods

**Participants** We recruited 449 participants (Age  $M = 30.83$ ;  $SD = 17.44$ ) on Prolific. They were randomly assigned to one of the three conditions of the experiment (Curiosity:  $N = 156$ ; Memory:  $N = 137$ ; Math:  $N = 156$ ). Participants were excluded if they showed irregular reaction times or their responses in the filler tasks indicates low engagement with the experiment. All exclusion criteria were pre-registered. The final sample included  $N$  participants (Curiosity  $N = 143$ ; Memory:  $N = 98$ ; Math:  $N = 139$ ).

**Procedure** This is a web-based self-paced visual presentation task. Participants were instructed to look at a sequence of animated creatures at their own pace and answer some questions throughout. At the end of the experiment, participants were asked to rate the similarity between pairs of creatures and complexity of creatures they encountered on a 7-point Likert Scale. Each participant saw eight blocks in total, half of which used creatures with high perceptual complexity, and half of which used creatures with low perceptual complexity. On each trial, an animated creature showed up on the screen. participants can press the down arrow to go to the next trial whenever they want after a minimum viewing time of 500 ms.

Each block consisted of six trials. A trial can be either a background trial (B) or a deviant trial (D). A background trial presented a creature repeatedly, and the deviant trial presented a different creature from the background trial in the block. Two creatures in the blocks were matched for visual complexity. There were four sequences of background trials and deviant trials. Each sequence appeared twice, once

with high complexity stimuli and once with low complexity stimuli. The deviant trial can appear at either the second (BDBBBB), the fourth (BBBDBB), or the sixth trial (BBBBBD) in the block. Two blocks do not have deviant trials (BBBBBB). The creatures presented in the deviant trials and background trials were matched for complexity. An overview of the experimental design can be seen in Figure X.




Figure 2: Experimental design and examples of simple and complex stimuli. In each block, a deviant could appear on the second, fourth (as depicted here) or sixth trial or not at all. Stimuli within a block were either all simple or all complex.

Participants were randomly assigned to one of the three conditions: Curiosity, Memory, and Math. The three conditions only differed in the type of questions asked following each block. In Curiosity condition, participants were asked to rate “How curious are you about the creature?” on a 5-point Likert scale. In Memory condition, a forced-choice recognition question followed each block (“Have you seen this creature before?”). The creature used in the question in both conditions was either a creature presented in the preceding block or a novel creature matched in complexity. In Math condition, the participants were asked a simple arithmetic question (“What is  $5 + 7$ ?”) in multiple-choice format.

**Stimuli** The animated creatures (Fig 1) were created using Spore (a game developed by Maxis in 2008). There were forty creatures in total, half of which have low perceptual complexity (e.g. the creatures do not have limbs, additional body parts, facial features, or textured skin), and half of which have high perceptual complexity (i.e. they do have the aforementioned features). We used the “animated avatar” function in Spore to capture the creatures in motion.

### Results

**Analytic Approach** The sample size, methods, and main analyses were all pre-registered and are available at [LINK]. Data and analysis scripts are available at [LINK].




Figure 3: Results of behavioral experiment.

**Manipulation Check** The complex animated creatures were rated as more perceptually complex ( $M = 5.2$ ;  $SD = 1.2$ ) than the simple animated creatures ( $M = 3.8$ ;  $SD = 1.1$ ). Pairs of background creature and deviant creature were rated as moderately dissimilar to one another ( $M = 3.5$ ;  $SD = 1.0$ ).

**Evaluating the Paradigm** Three criteria were selected to evaluate whether the paradigms successfully captured the characteristic looking time patterns observed in infant literature: habituation (the decrease in looking time for a stimulus with repeated presentations), dishabituation (the increase in looking time to a new stimulus after habituated to one stimulus), and complexity effect (longer looking time for perceptually more complex stimuli). To evaluate the phenomenon quantitatively, we ran a linear mixed effects model with maximal random effect structure. [DESCRIBE THE MODEL]. [REPORT THE MODEL RESULTS]

## Discussion

### Model comparison

#### Parameter estimation

We performed an iterative grid search in parameter space for each linking hypothesis. We a priori constrained our parameter space on the prior beta distribution to have shape parameters that  $\alpha_\theta > \beta_\theta$ , which describe the prior beliefs as “more likely to see the absence of a feature than the presence of a feature”. For model 1, we searched for the priors over the concept to be learned. The parameter search for the two metrics, surprisal and KL-divergence, converged on the same priors ( $\alpha_\theta = 1; \beta_\theta = 2$ ). For model 2, we searched for the priors over the concept ( $\theta$ ), the noise parameter that decides how likely a feature would be misperceived ( $\epsilon$ ), and the constant EIG from the world ( $EIG(world)$ ). The prior over the noise parameter was fixed for all searches ( $\alpha_\epsilon = 1; \beta_\epsilon =$

10). In model 2, different parameters were selected to obtain the best fit to the behavioral data (EIG:  $\alpha_\theta = 1, \beta_\theta = 4, \epsilon = 0.065, EIG(world) = 0.01$ ; KL:  $\alpha_\theta = 1, \beta_\theta = 5, \epsilon = 0.055, EIG(world) = 0.006$ ; Surprisal:  $\alpha_\theta = 1, \beta_\theta = 3, \epsilon = 0.07, EIG(world) = 8$ ).




Figure 4: Results from the discrete time model. The top panels show the trajectories of KL under different sequences, and the bottom panels show the trajectories of surprisal.

## Model experiment

To model the behavioral experiment, we first represented the stimuli as a vector of logical values indicating the presence and absence of a feature. All stimuli vectors are length 6, with the complex stimuli represented as having three TRUE and simple stimuli represented as having one TRUE. The rest of the elements are FALSE. Individual stimuli are then assembled into sequences to reflect the stimuli sequences in the behavioral experiment. For a particular sequence, we constructed the deviant stimulus based on the background stimulus to make sure that they were always maximally different and had the same number of features present.

For Model 1, since its behavior is non-probabilistic, we presented the model with each of the four sequences once and derived the information theoretic measurements. For model 2, we ran each sequence 500 times to obtain a reasonably precise estimate on the model’s behaviors.

## Results and Discussion

Both models reproduced the behavioral phenomena qualitatively, showing habituation, dishabituation, and complexity effect. To quantitatively explore the models, we fit the models’ output to the behavioral data. All models’ results were adjusted to match behavioral data’s scale and intercepts for easier comparisons. In model 1, we found that KL provided a better fit for the behavioral data than surprisal (Fig X; KL:

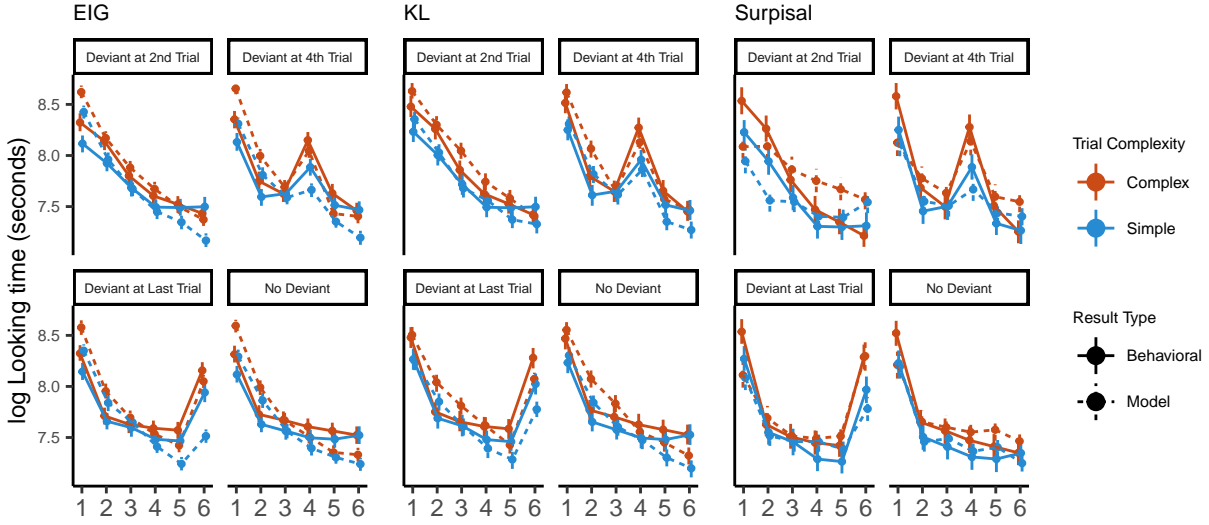


Figure 5: Continuous time model using different linking hypotheses provide qualitatively indistinguishable fits to the behavioral data. All model results are log-transformed and adjusted to be at the same scale and intercepts as the log-transformed behavioral data. The solid lines represent human data, and the dotted lines represent the model’s results. Red lines indicated results for complex stimuli, and blue lines indicated results for simple stimuli.

$r = 0.9$ ; RMSE = 0.39; Surprisal:  $r = 0.74$ ; RMSE = 0.42) In model 2, however, we found that the three linking hypotheses were qualitatively indistinguishable in their fits (Fig X; EIG:  $r = 0.92$ , RMSE = 0.19; KL:  $r = 0.93$ , RMSE = 0.12; Surprisal:  $r = 0.92$ , RMSE = 0.13). In this section, we discuss the implications of our findings and future directions.

First, the result from our discrete model extended the Poli et al (2020)’s findings. We showed in a new behavioral dataset that, similar to infants, adults’ attention allocation is better predicted by KL divergence than surprisal. This converging finding suggests a developmental continuity on principles guiding information-seeking behaviors. Second, the result from our continuous time model is the first evidence showing that under certain model architecture, surprisal and KL-divergence are good proxies for EIG, a metric that has the distinct advantage of quantitatively characterizing the optimal exploratory behaviors in humans. Tracking EIG can be computationally expensive and psychologically implausible. To calculate EIG, the current model needs to consider all possible combinations of features for the next observation. The proximities of model fits between EIG, KL, and surprisal revealed an opportunity to further optimize learning policies for computational models. Further, it also calls for more investigations on the relationships between different information theoretic metrics’ roles in predicting human behaviors. Our finding is in sharp contrast with previous work that tries to dissociate between different metrics using the same behavioral dataset (Poli, Emily). It would be theoretically interesting to consider under what circumstances do the measurements converge or become dissociable. Last but not least, although we can not directly compare the two models’ fits

quantitatively due to the differences in the number of free parameters, our results still show how models’ architectural differences can lead to different conclusions about the linking hypotheses. In a more psychologically realistic modeling regime, different information theoretic measurements can become indistinguishable.

There are several limitations to the current work. For our behavioral data, one concern is that our self-paced visual presentation task might not be capturing participants’ intrinsic interests in exploring the stimuli, which raises the question on whether we are actually measuring looking time. We addressed this question by including different filler tasks in-between blocks. No differences in looking time patterns are found across conditions. This suggests that the recorded-behaviors are task demand-independent, which means that they are more likely to reflect participants’ genuine interests in looking at the stimuli. For our models, there are several limitations that require further investigation. The current stimuli representation is rather oversimplified. For example, we did not take into consideration how features can have different degrees of saliency. In addition, the sampling policy’s implementation can be further challenged. The model currently decides between “continuing looking” and “look away”, but one can argue that in the behavioral experiment the participants were deciding between “continuing looking at the current stimulus” and “look at the next stimulus”. Building these more sophisticated assumptions into the model would certainly help us understand looking time better under a rational analysis framework. Nevertheless, our current work suggests that simpler models are capable of explaining key phenomena in looking time change.

Our ultimate goal is to provide a computational model that can explain the key looking time patterns documented and utilized in infant research: habituation, dishabituation, and complexity effect. We believe these patterns are driven by information-seeking principles that have strong developmental continuities. In the current work, we have shown that information theoretic measurements can be linked to adult looking time patterns. We also compared the two models and found that model architecture has consequences for linking hypotheses. Specifically, in a discrete time model, we found KL is a better predictor than surprisal. But in the continuous time model, three measurements (EIG, KL, and surprisal) become indistinguishable. As we further elaborate on our modeling approach, our ongoing work on infants will eventually help address the developmental trajectories of the mechanisms through which learners decide what to look at, and when to stop looking.

## References

- 10 Goodman, N. D., Tenenbaum, J. B., Feldman, J., & Griffiths, T. L. (2008). A rational analysis of rule-based concept learning. *Cognitive Science*, 32(1), 108–154.
- Kidd, C., Piantadosi, S. T., & Aslin, R. N. (2012). The goldilocks effect: Human infants allocate attention to visual sequences that are neither too simple nor too complex. *PloS One*, 7(5), e36399.
- Markant, D., & Gureckis, T. (2012). Does the utility of information influence sampling behavior? In *Proceedings of the annual meeting of the cognitive science society* (Vol. 34).
- Oaksford, M., & Chater, N. (1994). A rational analysis of the selection task as optimal data selection. *Psychological Review*, 101(4), 608.
- Poli, F., Serino, G., Mars, R., & Hunnius, S. (2020). Infants tailor their attention to maximize learning. *Science Advances*, 6(39), eabb5053.
- Tenenbaum, J. B. (1999). Bayesian modeling of human concept learning. *Advances in Neural Information Processing Systems*, 59–68.